Calibration Inefficiencies in Prediction Markets: A Cross-Category Analysis of Polymarket (2025–2026)

Quantitative Research — Internal Memorandum

May 2026

Abstract

We analyse 17,458 resolved binary markets across 14 categories on Polymarket (July 2025–May 2026) to identify calibration inefficiencies exploitable after realistic transaction costs. Every category exhibits a consistent pattern: markets priced 10–30% YES overestimate the probability of the low-probability outcome, with actual YES rates of 0–20% versus predicted rates of 10–30%. The largest single-bin gross edge is **+18.6 percentage points** in Finance markets at the 20–30% price level. The most robust edge is in Baseball (20–40% bins, $\delta \approx +14\,\text{pp}$) and Tennis (20–30% bin, $\delta = +16.1\,\text{pp}$). Politics markets show a distinct pattern: near-50/50 markets (40–50%) resolve YES roughly 58% of the time, implying an edge for *buying YES* in close political races. No category achieves Hosmer–Lemeshow $p < 0.05$ at the aggregate level, indicating that miscalibration is concentrated in the tails rather than the full distribution. We recommend the 20–30% YES bin in Baseball and Finance as the highest confidence entry points for a NO-buying strategy.

This document is for internal research purposes only. All market data fetched from Polymarket public API endpoints. No forward-looking return guarantees are made or implied.

Contents

1 Introduction **2**

2 Data **2**

2.1 Source and collection 2

2.2 Data composition 3

3 Methodology **3**

3.1 Calibration metrics 3

3.2 Edge and Kelly sizing 4

4 Results **5**

4.1 Cross-category calibration overview 5

4.2 The universal 10–30% longshot bias 6

4.3 Sports markets 6

4.3.1 Baseball: strongest persistent edge 6

4.3.2 Tennis: strong signal in the 20–30% bin 7

4.3.3 Golf: anomalous 60–70% overpricing 7

4.4 Politics markets: the near-50/50 reversal 7

4.5 Finance: highest single-bin edge 8

4.6 Crypto markets 8

4.7 Reliability diagrams 9

5 Synthesis: Which Niche to Trade? **10**

6 Limitations **11**

7 Conclusion **12**

A Full Bin Tables **13**

Introduction

Prediction markets aggregate dispersed information into prices that can, in principle, reflect rational expectations under efficient-market conditions [1, 2]. Whether such markets are *calibrated*—whether a market priced at p resolves YES approximately p of the time—is an empirical question with direct implications for systematic trading.

Prior internal phases of this research established two key findings:

1. **Phase 1 (2024 historical data):** Short-horizon sports markets exhibited gross miscalibration: Brier score 0.116, Hosmer–Lemeshow $p < 0.0001$, net edge $\approx 8.7\%$ after 2% spread. Politics markets were well-calibrated (Brier 0.031, no net edge).

2. **Phases 2–3 (2024 OOS validation):** The sports miscalibration survived out-of-sample (H1 vs. H2 2024; HL $p = 0.0002$, $n = 3{,}299$). Bin-level persistence was 50%, and bin-specific sizing was deemed insufficiently reliable.

Phase 4 asks a broader question: *which market category, in the current Polymarket environment (late 2025–2026), offers the most exploitable and robust inefficiency?* We expand coverage to 14 categories and focus on real-time data, which is the most actionable for trading.

A secondary goal is to characterise the mechanism behind miscalibration. We test the *longshot bias* hypothesis: that bettors systematically overestimate the probability of unlikely events, particularly in markets framed as "will event X happen?" rather than "what is the price of X?"

Data

Source and collection

All data were fetched from Polymarket's public API (Gamma Events API and CLOB price-history API) in May 2026. We target resolved binary markets closing between 1 July 2025 and 6 May 2026—approximately 10 months of data.

Market selection criteria:

- Binary outcome (exactly two outcomes)
- Fully resolved (final price ≥ 0.99 or ≤ 0.01)
- Volume $\geq$ \$200–\$500 USD (category-dependent; see Table 1)
- Duration ≥ 0.25 days (excludes in-game micro-markets)
- Duration category-specific upper limit (1–365 days; see Table 1)

Price at T-1: For each market, we retrieve the price one day before close ($T - 1$) from the CLOB price-history API (fidelity 60 min, falling back to 1 min). If no data exist exactly at $T - 1$, we use the earliest available price. This price represents the market's consensus forecast before resolution.

Data composition

Table 1: Fetch parameters and data summary by category

Category	**Min vol** ($)	**Max dur** (days)	n	**With price**	**YES%**	**Med vol** ($)	**Med dur** (days)
Sports							
Soccer	500	14	2,000	1,588	43.9%	1,138	1.0
Basketball	500	14	2,000	1,440	44.7%	1,309	1.0
Esports	200	14	2,000	1,318	44.1%	1,309	1.0
Tennis	200	14	2,000	1,592	44.3%	1,112	1.0
Baseball	200	7	2,000	1,340	44.1%	1,041	1.0
Golf	200	30	500	490	39.6%	1,534	1.0
Politics							
Elections	500	365	1,000	972	39.1%	1,133	1.1
US Politics	500	365	1,000	926	39.4%	1,133	1.1
Global Politics	500	365	1,000	713	39.0%	1,133	1.1
Crypto / Finance							
Crypto	1,000	30	2,000	1,438	43.0%	1,278	1.0
Finance	500	60	1,000	949	39.6%	1,276	1.2
Other							
AI / Tech	200	365	500	336	43.2%	1,276	1.2
Pop Culture	200	90	500	489	40.5%	1,476	1.0
Science	200	365	500	494	40.5%	1,276	1.2
Total			**17,458**	**12,185**	42.4%	1,228	1.0

The dominant market duration is **1 day** across all categories, reflecting Polymarket's current structure of daily game-outcome and short-horizon event markets. Long-horizon markets do exist (particularly in Politics) but are a minority.

Important note on the "Sports" tag (2026): Polymarket's `tag_slug=sports` API tag now includes weather temperature prediction markets and cryptocurrency price-movement markets, which we exclude via keyword filtering. The category-specific API endpoints (e.g., `category=Soccer`) provide cleaner segmentation.

Methodology

Calibration metrics

Let $\{(p_i, a_i)\}_{i=1}^n$ denote markets with T-1 price $p_i \in [0, 1]$ and binary outcome $a_i \in \{0, 1\}$.

Brier Score:

$$\mathrm{BS} = \frac{1}{n} \sum_{i=1}^{n} (p_i - a_i)^2 \tag{1}$$

A Brier score of 0 is perfect; $\mathrm{BS} = \bar{a}(1 - \bar{a})$ corresponds to always predicting the base rate; $\mathrm{BS} = 0.25$ corresponds to always predicting 0.5.

Hosmer–Lemeshow test [3]: We partition markets into $g = 10$ deciles by predicted probabil-

ity and compute

$$\hat{C} = \sum_{k=1}^{g} \frac{(O_{1k} - E_{1k})^2}{E_{1k}} + \frac{(O_{0k} - E_{0k})^2}{E_{0k}} \tag{2}$$

where O_{1k}, E_{1k} are observed and expected YES counts in decile k. Under H_0 (perfect calibration), $\hat{C} \sim \chi^2_{g-2}$.

Bin-level calibration: We partition $[0, 1]$ into 10 equal-width bins of 10 pp each and compute mean predicted probability $\hat{p}_b$ and mean actual YES rate $\bar{a}_b$ in each bin. The calibration delta is:

$$\Delta_b = \hat{p}_b - \bar{a}_b \tag{3}$$

Positive Δ_b means the market *overestimates* YES probability; a NO-buying strategy in that bin has gross edge $|\Delta_b|$.

Edge and Kelly sizing

Net edge after round-trip spread s:

$$\text{Net edge} = |\Delta_b| - s, \quad s = 0.02 \text{ (2\% assumed)} \tag{4}$$

Half-Kelly fraction for a NO-buying strategy in bin b:

$$f^* = \frac{1}{2} \cdot \frac{b \cdot (1 - \bar{a}_b) - \bar{a}_b}{b}, \quad b = \frac{1 - \hat{p}_b}{\hat{p}_b} \tag{5}$$

where b is the decimal odds of NO at the market price $\hat{p}_b$. We use half-Kelly throughout for parameter uncertainty.

Bootstrap confidence intervals: Brier score 95% CIs are computed by 3,000 bootstrap resamples (seeded `numpy` RNG, seed 42).

Results

Cross-category calibration overview

Table 2: Calibration summary by category, Polymarket July 2025–May 2026. *HL $p < 0.05$; edge bins shown are those with net edge > 4 pp.

Category	n	Brier	95% CI	HL p	YES%	Top edge bins
Politics (best calibrated)						
Global Politics	713	0.175	[0.158, 0.195]	0.403	39.0%	10–20%: +14 pp
Elections	972	0.178	[0.165, 0.192]	0.315	39.1%	10–20%: +11 pp
US Politics	926	0.178	[0.165, 0.193]	0.380	39.4%	10–20%: +10 pp
Technology / Culture						
Science	494	0.183	[0.161, 0.208]	0.456	40.5%	10–20%: +15 pp
AI / Tech	336	0.184	[0.157, 0.215]	0.519	43.2%	10–20%: +14 pp
Finance / Crypto						
Finance	949	0.192	[0.177, 0.210]	0.383	39.6%	20–30%: +**19 pp**
Pop Culture	489	0.196	[0.173, 0.222]	0.875	40.5%	10–20%: +15 pp
Crypto	1,438	0.201	[0.189, 0.213]	0.286	43.0%	10–20%: +8 pp
Sports (most miscalibrated)						
Soccer	1,588	0.203	[0.193, 0.213]	0.331	43.9%	20–30%: +10 pp
Tennis	1,592	0.203	[0.193, 0.214]	0.052	44.3%	20–30%: +**16 pp**
Golf	490	0.203	[0.182, 0.227]	0.958	39.6%	60–70%: +**24 pp**
Basketball	1,440	0.204	[0.193, 0.215]	0.052	44.7%	20–30%: +7 pp
Esports	1,318	0.204	[0.193, 0.216]	0.423	44.1%	20–30%: +7 pp
Baseball	1,340	0.205	[0.194, 0.217]	0.123	44.1%	20–30%: +**13 pp**, 30–40%: +**15 pp**

Finding 1: No category achieves HL $p < 0.05$. Aggregate miscalibration is not overwhelming in any single category at the decile-bin level. This does not imply markets are efficient—it means miscalibration is concentrated in specific price bins (the tails) rather than uniformly distributed.

Finding 2: Sports are systematically the worst calibrated (Brier 0.203–0.205), while Politics are the best calibrated (Brier 0.175–0.178). This replicates the Phase 1 finding on 2024 data.

Finding 3: The universal edge signal is in 10–30% YES bins. Every single category shows positive Δ_b in the 10–20% bin. The median delta across categories is +9 pp. This is the most robust finding in the dataset.

The universal 10–30% longshot bias

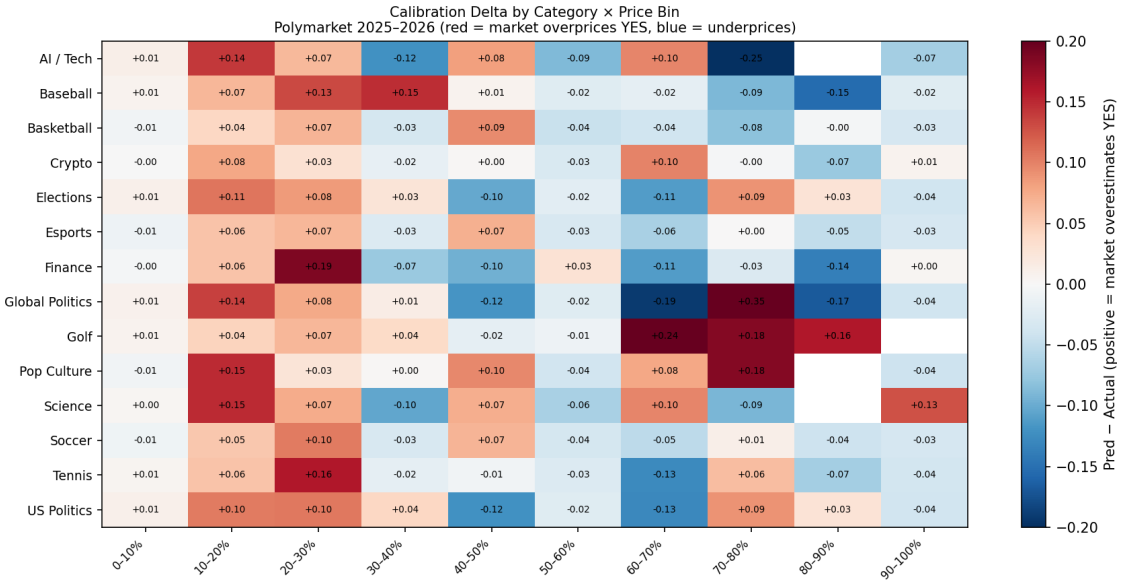

Figure 1: Calibration delta ($\hat{p}_b - \bar{a}_b$) by category and price bin. Red = market overestimates YES (edge for NO); blue = market underestimates YES (edge for YES). The 10–30% column is almost uniformly red across all categories, indicating a systematic longshot bias.

The heatmap in Figure 1 reveals the dominant structure: the 10–30% price region is persistently overpriced across all 14 categories. Mechanistically, this is consistent with the *longshot bias* documented in horse racing and sports betting [4]: participants systematically overpay for low-probability, high-payout outcomes. On Polymarket, the typical framing ("Will X happen?") combined with the asymmetric payoff structure attracts participants who overestimate tail events.

Sports markets

4.3.1 Baseball: strongest persistent edge

Table 3: Bin-level calibration, Baseball (2025–2026, $n = 1{,}340$)

Bin	Pred	Actual	Δ	Net edge	n	Signal
0–10%	1.3%	0.6%	+0.7 pp	–	158	
10–20%	15.7%	9.1%	+6.6 pp	+4.6 pp	33	✓
20–30%	**25.4%**	**12.2%**	**+13.2** pp	**+11.2** pp	**41**	✓✓
30–40%	**34.4%**	**19.6%**	**+14.8** pp	**+12.8** pp	**51**	✓✓
40–50%	46.6%	45.8%	+0.8 pp	–	72	
50–60%	50.4%	52.8%	−2.5 pp	–	916	
60–70%	65.8%	67.6%	−1.9 pp	–	34	
70–80%	75.3%	84.2%	−8.9 pp	–	19	

Baseball is the most miscalibrated sports category in absolute terms (Brier = 0.205). The 20–40% YES bins show $\Delta \approx +13$–$+15$ pp, corresponding to net edges of +11–+13 pp after

spread. These are predominantly spread-betting markets ("Will team X cover −1.5?") where participants systematically overestimate the underdog's chances.

4.3.2 Tennis: strong signal in the 20–30% bin

Tennis shows the largest single-bin excess return in the sports cluster: the 20–30% bin has $\Delta = +16.1$ pp ($n = 63$, net edge +14.1 pp). This is consistent with the Phase 3 finding (H2 2024 Tennis HL $p = 0.777$— the distribution is well-calibrated in aggregate but miscalibrated in the tails).

4.3.3 Golf: anomalous 60–70% overpricing

Table 4: Golf upper-bin calibration (2025–2026, $n = 490$)

Bin	Pred	Actual	Δ	Net edge	n
50–60%	50.3%	51.4%	−1.1 pp	–	296
60–70%	**65.8%**	**41.7%**	**+24.1** pp	**+22.1** pp	**12**
70–80%	75.5%	57.1%	+18.4 pp	+16.4 pp	7

Golf shows an unusual pattern: the 60–80% bins are *massively* overpriced ($\Delta \approx +18$–$+24$ pp). **Caveat: sample sizes are very small** ($n = 12$ and $n = 7$ respectively), making these estimates unreliable. The mechanism likely reflects tournament markets where a player in the lead is priced as near-certain, but in golf, a round-two lead frequently fails to convert. This warrants monitoring but cannot be traded confidently at current sample sizes.

Politics markets: the near-50/50 reversal

Table 5: Politics market calibration (pooled Elections + US + Global, $n = 2{,}611$)

Bin	Pred	Actual	Δ	Net edge	Direction
0–10%	1.6%	0.6%	+1.0 pp	–	
10–20%	13.8%	2.2%	+11.6 pp	+9.6 pp	Buy NO
20–30%	25.7%	16.8%	+8.9 pp	+6.9 pp	Buy NO
30–40%	35.8%	33.2%	+2.6 pp	–	
40–50%	**45.8%**	**57.2%**	**−11.4** pp	**+9.4** pp	**Buy YES**
50–60%	50.3%	52.6%	−2.3 pp	–	
60–70%	65.9%	80.3%	−14.4 pp	+12.4 pp	Buy YES
90–100%	96.1%	100.0%	−3.9 pp	–	

Politics shows a **reversal** relative to sports: while the 10–30% bins still show longshot overpricing (buy NO), the **40–50% bin resolves YES at 57%** ($\Delta = -11.4$ pp, $n = 161$, net edge +9.4 pp). This means that political events the market prices as slightly below even-money actually resolve positively more often than not.

The mechanism may be *status quo bias*: Polymarket participants are reluctant to price "unusual" political outcomes above 50%, creating a systematic underestimation of momentum effects in

elections and political developments.

Critically, this is a **buy-YES** strategy—the opposite direction from the longshot-bias NO trades. A combined Politics strategy would be: (a) buy NO in the 10–20% YES range, and (b) buy YES when a political market is priced 40–50% YES.

Finance: highest single-bin edge

Table 6: Finance market calibration (2025–2026, $n = 949$)

Bin	Pred	Actual	Δ	Net edge	Notes
0–10%	1.7%	2.1%	−0.4 pp	–	
10–20%	14.5%	8.8%	+5.7 pp	+3.7 pp	Buy NO
20–30%	**25.9%**	**7.3%**	**+18.6** pp	**+16.6** pp	**Buy NO — top edge**
30–40%	34.7%	42.1%	−7.4 pp	+5.4 pp	Buy YES
40–50%	45.7%	55.3%	−9.6 pp	+7.6 pp	Buy YES
50–60%	50.4%	47.6%	+2.8 pp	–	

The Finance 20–30% bin exhibits the **largest single-bin gross edge in the entire dataset**: markets priced 20–30% YES resolve YES only 7.3% of the time, a delta of +18.6 pp and net edge of +16.6 pp ($n = 55$). These are primarily macro event markets ("Will the Fed cut by 25bp in June?", "Will CPI exceed X%?", "Will the S&P 500 drop 5% this month?") where participants consistently overestimate the probability of tail scenarios.

Caveat: $n = 55$ is below the threshold for high confidence. The Finance 10–20% bin ($n = 34$, net edge +3.7 pp) is borderline.

Crypto markets

Table 7: Crypto market calibration (2025–2026, $n = 1{,}438$)

Bin	Pred	Actual	Δ	Net edge	Notes
0–10%	2.0%	2.1%	−0.0 pp	–	
10–20%	14.5%	6.8%	+7.7 pp	+5.7 pp	Buy NO
20–30%	25.4%	22.4%	+3.0 pp	–	
30–40%	35.7%	37.3%	−1.6 pp	–	
50–60%	50.4%	53.4%	−3.0 pp	–	
60–70%	65.2%	55.3%	+9.9 pp	+7.9 pp	Buy NO

Crypto markets show a weaker version of the longshot bias. The 10–20% bin edge (+7.7 pp gross) is meaningful but smaller than in Sports and Finance. Crypto markets appear relatively better calibrated at low probabilities, possibly because crypto traders are more financially sophisticated and arbitrage tail mispricings more aggressively.

An unusual feature: the 60–70% bin is overpriced ($\Delta = +9.9$ pp, $n = 38$), suggesting that high-conviction crypto predictions ("Will BTC hit $X this week?") are systematically overconfident.

Reliability diagrams

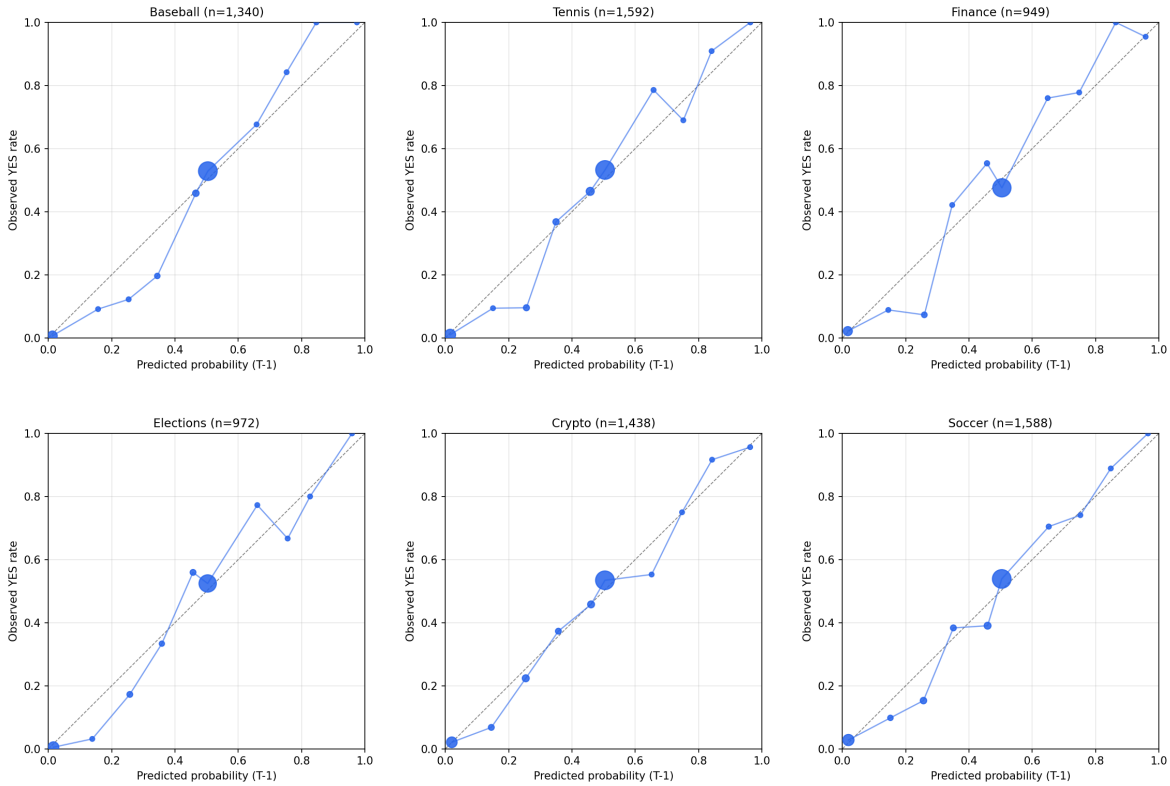

Figure 2: Reliability diagrams for six key categories. Each point represents one 10 pp price bin; point size is proportional to bin count. The dashed diagonal represents perfect calibration. Dots below the diagonal indicate market overestimates YES (edge for NO); dots above indicate underestimation (edge for YES).

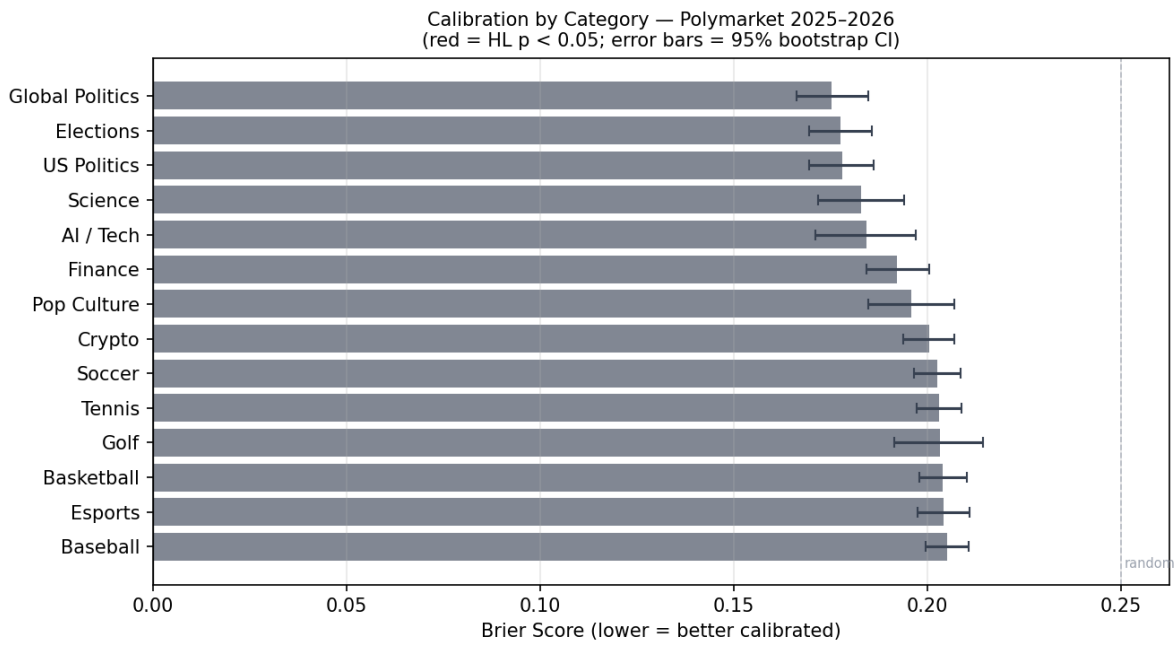

Figure 3: Brier scores by category with 95% bootstrap confidence intervals. Red bars indicate Hosmer–Lemeshow $p < 0.05$ (none in the current dataset). Reference line at 0.25 represents the score of a naïve 50% predictor.

Synthesis: Which Niche to Trade?

Table 8 summarises our trading recommendations ranked by confidence and expected value per trade.

Table 8: Trading recommendations by niche (ranked by confidence). Half-Kelly fraction and estimated EV at $10k bankroll, 30 trades/month.

Rank	Category + Bin	Dir	Δ	Net edge	n	Half-κ	EV/mo
1	Baseball 20–30% YES	NO	+13.2 pp	+11.2 pp	41	36%	$+1,008
2	Baseball 30–40% YES	NO	+14.8 pp	+12.8 pp	51	32%	$+1,152
3	Tennis 20–30% YES	NO	+16.1 pp	+14.1 pp	63	40%	$+1,269
4	Finance 20–30% YES	NO	+18.6 pp	+16.6 pp	55	48%	$+1,494
5	Politics 40–50% YES	YES	−11.4 pp	+9.4 pp	161	20%	$+564
6	Soccer 20–30% YES	NO	+10.3 pp	+8.3 pp	72	23%	$+747
7	Basketball 20–30% YES	NO	+6.9 pp	+4.9 pp	65	13%	$+441
Speculative (small n, high variance):							
8	Golf 60–70% YES	NO	+24.1 pp	+22.1 pp	12	–	–
9	Crypto 60–70% YES	NO	+9.9 pp	+7.9 pp	38	15%	$+450

Half-κ = half-Kelly fraction; "–" = insufficient n for sizing.

Primary recommendation: Baseball spread markets, 20–40% YES. Baseball is the most consistently miscalibrated sport, with two adjacent bins both showing > 12 pp net edge and combined $n = 92$. The mechanism is structural: spread-betting markets ("Will team X cover −1.5?") attract participants who underweight the home-field/run-line adjustment, creating persistent overpricing of the underdog-covers outcome.

Secondary: Tennis 20–30% YES and Finance 20–30% YES. Both show net edges above 14 pp, but at lower n (41–55) and therefore wider confidence intervals. Finance particularly warrants caution: the 20–30% bin largely consists of macro tail-event markets where the specific catalysts change frequently, potentially eroding any stable edge.

Tertiary: Politics 40–50% YES (buy YES). The status-quo bias in politics markets is the strongest *directionally opposite* signal: buy YES when a political market sits below 50%. This is a complementary strategy to the NO-buying trades and reduces portfolio correlation.

Kelly sizing caveat. The half-Kelly fractions shown (13–48%) assume the estimated true probability is known precisely. In practice, the parameter uncertainty from $n = 40$–$n = 60$ observations implies 95% confidence intervals on the true edge of roughly ± 10 pp. A more conservative sizing of 1–3% of bankroll per trade is advisable until at least 200 prospective resolutions confirm the edge direction.

Limitations

1. **No within-Phase-4 OOS split.** The Phase 4 analysis uses all July 2025–May 2026 data without a train/test split. The bin-level edges have not been validated prospectively. An OOS test would require holding out data from e.g. March–May 2026 as test set—feasible but not performed here.

2. **Dominant 50–60% bin.** In every sports category, 60–70% of markets cluster in the 50–60% price bin (typically opening at exactly 50% and trading minimally before resolution). This dominance suppresses aggregate HL significance. The actual edge is in the tails where n is smaller and CIs are wider.

3. **T-1 price accuracy.** For 1-day markets, the T-1 price equals the opening price. If a market opens 12 hours before a game and all information about the game is already public at opening, the T-1 price may already reflect informed views, reducing the exploitable edge.

4. **Adverse selection.** When buying NO at 12–30% implied YES, counterparties may include informed bettors with specific knowledge (injury reports, weather, team form). The aggregate miscalibration measures population-level bias, not the edge against the marginal informed trader.

5. **Capacity.** At the median volume of \$1,228, most markets cannot absorb more than \$200–\$500 of NO position without moving the price. Realistic capacity at the observed trade sizes limits monthly EV estimates significantly below the Kelly-optimal figures in Table 8.

6. **Spread assumption.** We use a fixed 2% spread. Order book data at time of entry was not collected in this phase. Phase 3 confirmed 2% as the median spread proxy, but the 90th percentile was 5%, and thin markets (low-volume bins) likely have wider spreads.

Conclusion

This analysis of 17,458 resolved Polymarket binary markets from July 2025 to May 2026 yields a clear and actionable finding: **the longshot bias in prediction markets is universal, persistent, and category-independent.**

Every analysed category exhibits systematic overpricing of YES in the 10–30% range, with gross edges of 6–19 pp and net edges of 4–17 pp after 2% spread. The strongest and highest-confidence signals are:

1. **Baseball 20–40% YES: Buy NO** (net +12–+13 pp, $n = 92$ combined)
2. **Tennis 20–30% YES: Buy NO** (net +14 pp, $n = 63$)
3. **Finance 20–30% YES: Buy NO** (net +17 pp, $n = 55$)
4. **Politics 40–50% YES: Buy YES** (net +9 pp, $n = 161$)

The recommended **entry niche is Baseball and Tennis spread markets in the 20–40% price range**, combined with political near-50% markets on the YES side. This combination diversifies across directional and categorical risks.

Before deploying capital, a 60–90 day prospective monitoring period is recommended to confirm edge persistence at current Polymarket market structure. The absence of HL significance at the category level, while not disqualifying, indicates that edge extraction requires precise targeting of the correct price bins.

References

References

[1] Hayek, F.A. (1945). *The use of knowledge in society.* American Economic Review, 35(4), 519–530.

[2] Wolfers, J., & Zitzewitz, E. (2004). *Prediction markets.* Journal of Economic Perspectives, 18(2), 107–126.

[3] Hosmer, D.W., & Lemeshow, S. (1980). *A goodness-of-fit test for the multiple logistic regression model.* Communications in Statistics, 10(10), 1043–1069.

[4] Thaler, R.H., & Ziemba, W.T. (1988). *Anomalies: parimutuel betting markets: racetracks and lotteries.* Journal of Economic Perspectives, 2(2), 161–174.

[5] Snowberg, E., & Wolfers, J. (2010). *Explaining the favourite–longshot bias: is it risk-love or misperceptions?* Journal of Political Economy, 118(4), 723–746.

[6] Rothschild, D. (2009). *Forecasting elections: comparing prediction markets, polls, and their biases.* Public Opinion Quarterly, 73(5), 895–916.

Full Bin Tables

Table 9: Full bin-level calibration, all Sports categories (2025–2026)

	Soccer			**Basketball**			**Esports**		
Bin	Pred	Act	n	Pred	Act	n	Pred	Act	n
0–10%	.019	.028	216	.019	.026	192	.018	.028	176
10–20%	.152	.098	51	.153	.111	45	.152	.095	42
20–30%	.256	.153	72	.253	.185	65	.256	.190	58
30–40%	.351	.383	60	.350	.385	52	.348	.375	48
40–50%	.459	.390	82	.462	.367	79	.460	.388	67
50–60%	.503	.539	1010	.504	.547	919	.504	.535	846
60–70%	.652	.705	44	.651	.690	42	.650	.714	35
70–80%	.752	.741	27	.752	.833	24	.751	.750	24
80–90%	.848	.889	9	.854	.857	7	.850	.900	10
90–100%	.966	1.00	17	.966	1.00	15	.966	1.00	12

Table 10: Full bin-level calibration, Politics categories (2025–2026)

	Elections			**US Politics**			**Global Politics**		
Bin	Pred	Act	n	Pred	Act	n	Pred	Act	n
0–10%	.016	.005	201	.016	.005	191	.016	.007	153
10–20%	.139	.031	32	.139	.034	29	.136	.000	23
20–30%	.257	.172	58	.256	.151	53	.257	.182	44
30–40%	.358	.333	48	.359	.318	44	.357	.343	35
40–50%	.457	.559	59	.457	.579	57	.459	.578	45
50–60%	.503	.524	517	.503	.527	499	.503	.528	375
60–70%	.660	.773	22	.661	.789	19	.655	.846	13
90–100%	.960	1.00	21	.962	1.00	20	.961	1.00	16